

3.1 WHAT IS `autograd`

01 September 2025

12:37 AM

👤 Avinash Yadav

`autograd` is a core component of PyTorch that *provides automatic differentiation for tensor operations*. It *enables gradient computation*, which is essential for training machine learning models using optimization algorithms like gradient descent.

↳ With the help of it, we can perform automatic differentiation on any kind of tensor's operation

When training neural networks, the most frequently used algorithm is **back propagation**. In this algorithm, parameters (model weights) are adjusted according to the **gradient** of the loss function with respect to the given parameter. To compute those gradients, PyTorch has a built-in differentiation engine called `torch.autograd`. It supports automatic computation of gradient for any computational graph.

EXAMPLE 1.1

$$1.1 \quad y = x^2 \xrightarrow{\text{find}} \frac{dy}{dx} \quad \text{for given 'x'}$$

```

1 import torch
✓ 0.0s

```

Scalar Tensor

```

1 x = torch.tensor(3.0, requires_grad=True)
✓ 0.0s

```

we are setting requires_grad attribute as True. by default it is False; By this we are telling PyTorch that we want to calculate derivative of tensor 'x'. so it starts tracking it.

```

1 y = x**2
✓ 0.0s

```

y is calculated in forward pass

```

1 print(x)
2 print(y)
✓ 0.0s

```

So when we set it True, PyTorch make computation graph internally

Forward: $y = x^2$

Backward: $\frac{dy}{dx}$

tensor(3., requires_grad=True)
tensor(9., grad_fn=<PowBackward0>)

computed value

```

1 y.backward()
✓ 0.0s

```

now going backward, it will compute $\frac{dy}{dx}$ i.e. $2x$

```

1 x.grad
✓ 0.0s

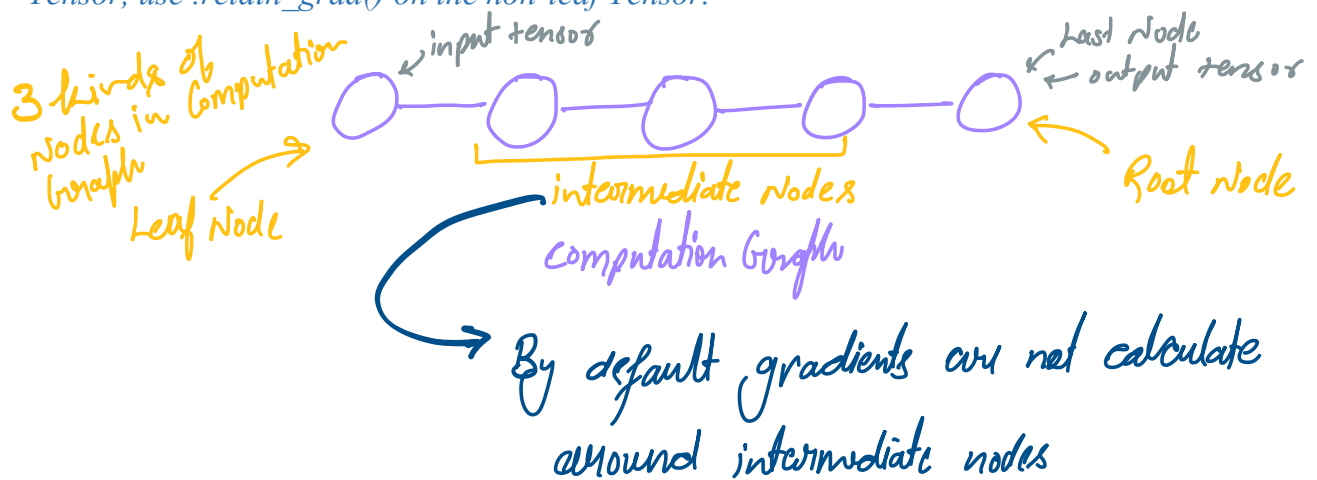
```

this is $\frac{dy}{dx} = 2x = 2(3) = 6$

tensor(6.)

So while going backward PyTorch needs to know what operation was performed here while it went for forward pass. that's why it remembers it as 'PowBackward0' for backward pass

The `.grad` attribute of a Tensor that is not a leaf Tensor won't be populated during `autograd.backward()`. If you indeed want the `.grad` field to be populated for a non-leaf Tensor, use `.retain_grad()` on the non-leaf Tensor.



Conceptually, autograd keeps a record of data (tensors) and all executed operations (along

with the resulting new tensors) in a directed acyclic graph (DAG) consisting of Function objects. In this DAG, *leaves are the input tensors, roots are the output tensors*. By tracing this graph from roots to leaves, you can automatically compute the gradients using the chain rule.

In a **forward pass**, autograd does two things simultaneously:

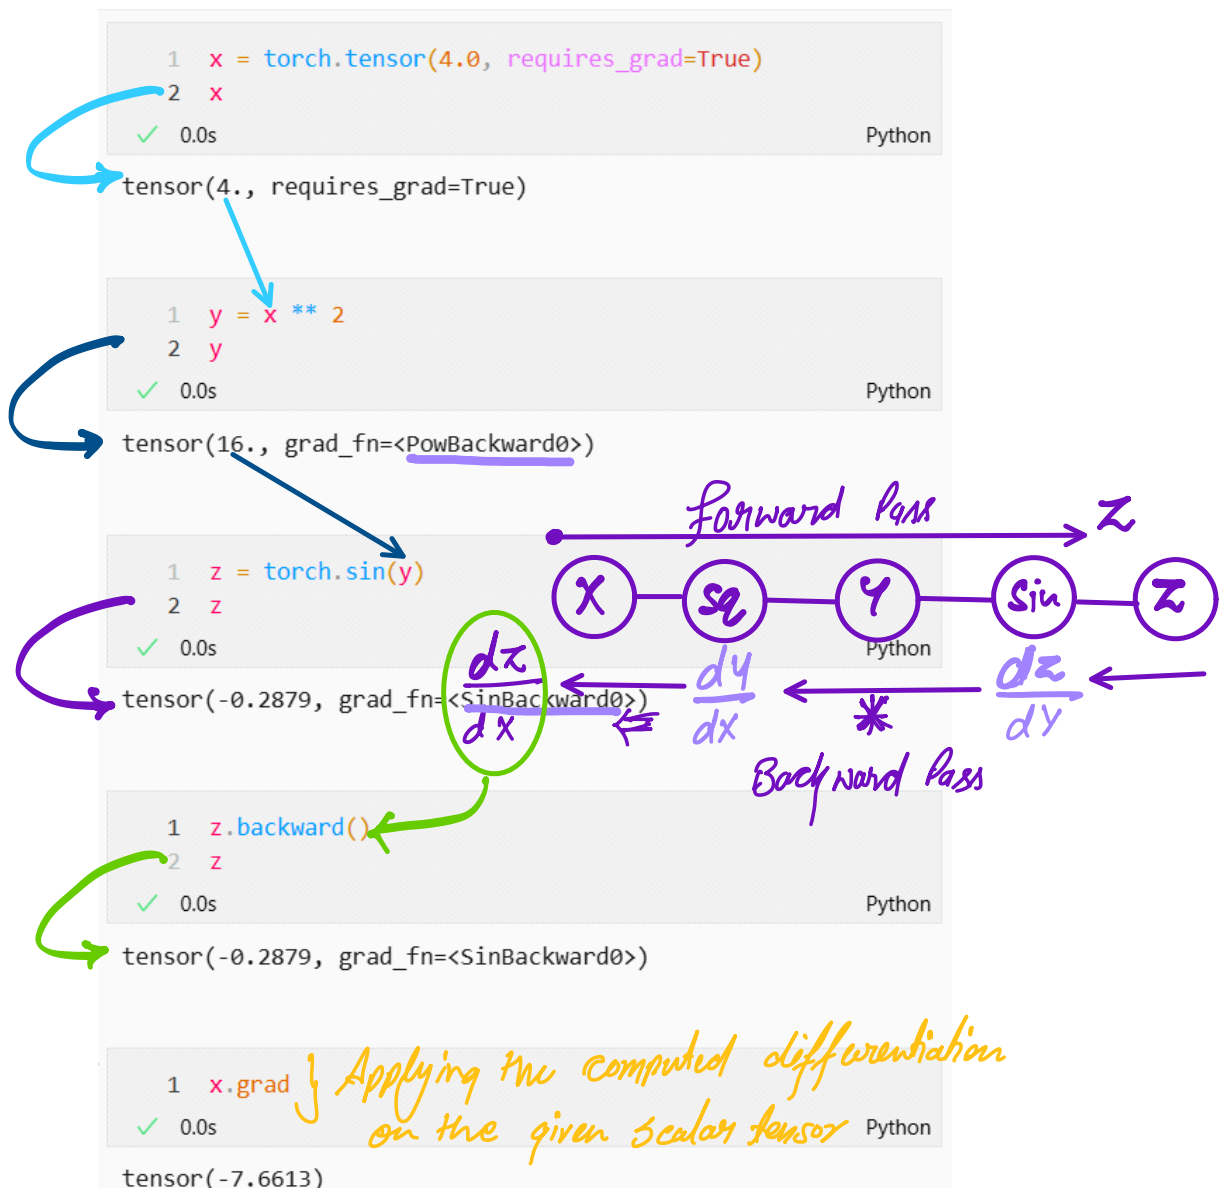
- run the requested operation to compute a resulting tensor
- maintain the operation's *gradient function* in the DAG.

The **backward pass** kicks off when *.backward()* is called on the DAG *root.autograd* then:

- computes the gradients from each *.grad_fn*,
- accumulates them in the respective tensor's *.grad attribute*
- using the chain rule, propagates all the way to the leaf tensors.

EXAMPLE 2.1

$$\begin{array}{l} y = x^2 \\ z = \sin(y) \end{array} \quad \left. \vphantom{\begin{array}{l} y = x^2 \\ z = \sin(y) \end{array}} \right\} \text{ we need } \frac{dz}{dx}$$

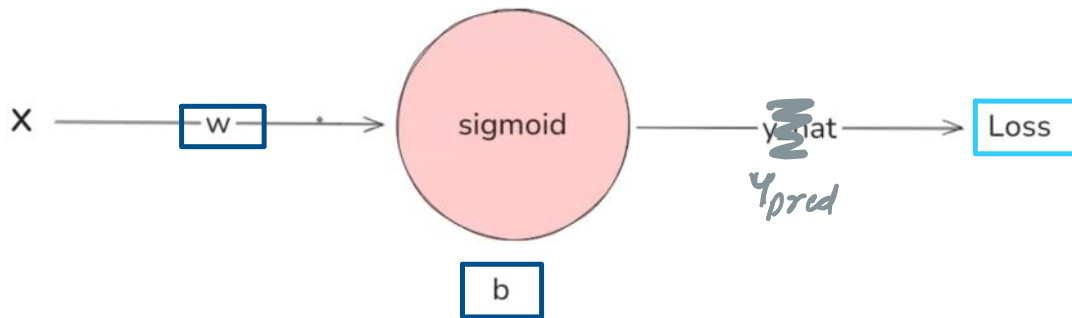


How Does autograd works?

Input



CGPA	PLACED
6.7	0 \rightarrow No



1. Linear Transformation:

$$z = w \cdot x + b$$

2. Activation (Sigmoid Function):

$$y_{pred} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3. Loss Function (Binary Cross-Entropy Loss):

$$L = -[y_{target} \cdot \ln(y_{pred}) + (1 - y_{target}) \cdot \ln(1 - y_{pred})]$$

since we have 2 parameters w & b . Therefore we'll need to find 2 derivatives.

1st derivative of loss wrt $w \Rightarrow \frac{\partial L}{\partial w}$
 2nd derivative of loss wrt $b \Rightarrow \frac{\partial L}{\partial b}$

$$\therefore \frac{\partial L}{\partial w} = \underbrace{\frac{\partial L}{\partial y_{pred}}}_{\text{from eq. 3}} \times \underbrace{\frac{\partial y_{pred}}{\partial z}}_{\text{from eq. 2}} \times \underbrace{\frac{\partial z}{\partial w}}_{\text{from eq. 1}}$$

Similarly;

$$\frac{\partial L}{\partial b} = \underbrace{\frac{\partial L}{\partial y_{pred}}}_{\text{from eq. 3}} \times \underbrace{\frac{\partial y_{pred}}{\partial z}}_{\text{from eq. 2}} \times \underbrace{\frac{\partial z}{\partial b}}_{\text{from eq. 1}}$$

So now computing derivatives:

$$\frac{\partial L}{\partial y_{pred}} = \frac{(y_{pred} - y)}{y_{pred}(1 - y_{pred})}$$

$$\frac{\partial z}{\partial w} = x$$

$$\frac{\partial y_{pred}}{\partial z} = y_{pred}(1 - y_{pred})$$

$$\frac{\partial z}{\partial b} = 1$$

$$\therefore \frac{\partial L}{\partial w} = \frac{\partial L}{\partial y_{pred}} \times \frac{\partial y_{pred}}{\partial z} \times \frac{\partial z}{\partial w}$$

$$= \frac{(y_{pred} - y)}{y_{pred}(1 - y_{pred})} \times y_{pred}(1 - y_{pred}) \times x$$

$$= (y_{pred} - y) \times x$$

$$\therefore \frac{\partial L}{\partial w} = x(y_{pred} - y)$$

Similarly

$$\frac{\partial L}{\partial b} = (y_{pred} - y)$$

```
1 import torch
```

```
2
```

```
3 # Inputs
```

```
4 x = torch.tensor(6.7) # Input feature
```

→ x

```
5 y = torch.tensor(0.0) # True Label (binary)
```

→ y

```
6
```

```
7 w = torch.tensor(1.0) # Weight
```

y initial setup with Random value

```
8 b = torch.tensor(0.0) # Bias
```

✓ 0.0s

Python

$$L = -[y_{\text{target}} \cdot \ln(y_{\text{pred}}) + (1 - y_{\text{target}}) \cdot \ln(1 - y_{\text{pred}})]$$

Calculating
loss

```
1 # Binary Cross-Entropy Loss for scalar
```

```
2 def binary_cross_entropy_loss(prediction, target):
```

```
3     epsilon = 1e-8 # To prevent log(0)
```

```
4     prediction = torch.clamp(prediction, epsilon, 1 - epsilon)
```

```
5     return -(target * torch.log(prediction) + (1 - target) * torch.log(1 - prediction))
```

✓ 0.0s

Python

```
1 # Forward pass
```

```
2 z = w * x + b # Weighted sum (linear part)
```

→ Eq. 1

```
3
```

```
4 y_pred = torch.sigmoid(z) # Predicted probability
```

→ Eq. 2

```
5
```

```
6 # Compute binary cross-entropy loss
```

```
7 loss = binary_cross_entropy_loss(y_pred, y)
```

→ Eq. 3

```
8 loss
```

✓ 0.0s

Python

tensor(6.7012)

BACK-PROPAGATION

```
1 # Derivatives:
```

```
2 # 1. dL/d(y_pred): Loss with respect to the prediction (y_pred)
```

```
3 dloss_dy_pred = (y_pred - y)/(y_pred*(1-y_pred))
```

```
4
```

```
5 # 2. dy_pred/dz: Prediction (y_pred) with respect to z (sigmoid derivative)
```

```
6 dy_pred_dz = y_pred * (1 - y_pred)
```

```
7
```

```
8 # 3. dz/dw and dz/db: z with respect to w and b
```

```
9 dz_dw = x # dz/dw = x
```

```
10 dz_db = 1 # dz/db = 1 (bias contributes directly to z)
```

```
11
```

```
12 dL_dw = dloss_dy_pred * dy_pred_dz * dz_dw
```

```
13 dL_db = dloss_dy_pred * dy_pred_dz * dz_db
```

✓ 0.0s

Python

```
1 print(f"Manual Gradient of loss w.r.t weight (dw): {dL_dw}")
```

```
2 print(f"Manual Gradient of loss w.r.t bias (db): {dL_db}")
```

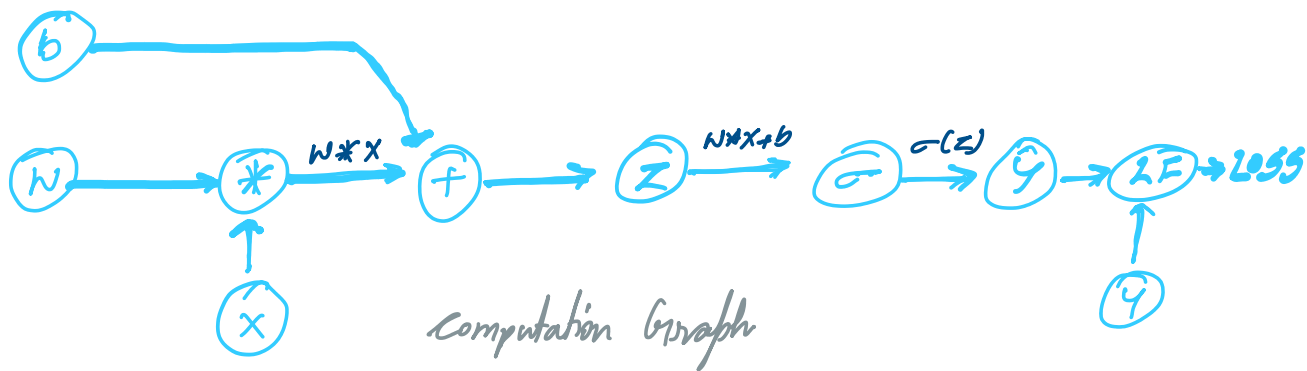
✓ 0.0s

Python

Manual Gradient of loss w.r.t weight (dw): 6.691762447357178

Manual Gradient of loss w.r.t bias (db): 0.998770534992218

— x — x — x — x — x — x —



Using autograd

```
1 x = torch.tensor(6.7)
2 y = torch.tensor(0.0)
3 print(x)
4 print(y)
```

✓ 0.0s

Python

```
tensor(6.7000)
tensor(0.)
```

we are not using 'requires_grad=True' because we do not have to calculate the derivatives w.r.t to x & y

```
1 w = torch.tensor(1.0, requires_grad=True)
2 b = torch.tensor(0.0, requires_grad=True)
3 print(w)
4 print(b)
```

✓ 0.0s

Python

```
tensor(1., requires_grad=True)
tensor(0., requires_grad=True)
```

STARTING FORWARD PASS:

```
1 z = w*x + b
2 z
```

Eq. 1

✓ 0.0s

Python

```
tensor(6.7000, grad_fn=<AddBackward0>)
```

```
1 y_pred = torch.sigmoid(z)
2 y_pred
```

Eq. 2

✓ 0.0s

Python

```
tensor(0.9988, grad_fn=<SigmoidBackward0>)
```

```
1 loss = binary_cross_entropy_loss(y_pred, y)
2 loss
```

Eq. 3

✓ 0.0s

Python

```
tensor(6.7012, grad_fn=<NegBackward0>)
```

```
1 loss.backward()
```

✓ 0.0s

Python


```
1 loss.backward()
```

✓ 0.0s

Python

```
1 print(f"By using autograd function (dw): {w.grad}")
2 print(f"By using autograd function (db): {b.grad}")
```

✓ 0.0s

Python

By using autograd function (dw): 6.6917619705200195

By using autograd function (db): 0.9987704753875732



Clearing Gradients

The key concept behind `autograd` in PyTorch is **gradient accumulation**. Each time you call `.backward()` on a tensor, the gradients for all tensors with `requires_grad=True` are added (accumulated) to their `.grad` attributes. This means if you run the backward pass multiple times without clearing the gradients, the values in `.grad` will keep increasing, reflecting the sum of all computed gradients.

Why does this happen?

This behavior is useful when training neural networks using mini-batches. You might want to accumulate gradients over several batches before updating the model parameters.

How to manage gradient accumulation:

- To avoid unwanted accumulation, always clear gradients before a new backward pass using `.zero_()` on the `.grad` attribute:

```
x.grad.zero_()
```

- Alternatively, use `optimizer.zero_grad()` when working with optimizers.

Summary:

- `.backward()` accumulates gradients in `.grad`.
- Always clear gradients before a new backward pass unless you intentionally want to accumulate them.

FIRST RUN

```
1 x = torch.tensor(2.0, requires_grad=True)
2 x
```

✓ 0.0s

Python

tensor(2., requires_grad=True)

SECOND RUN

FP

```
1 y = x ** 2
2 y
```

✓ 0.0s

FP

```
1 y = x ** 2
2 y
```

✓ 0.0s

FP

```
1 y = x ** 2
2 y
```

✓ 0.0s

THIRD RUN

tensor(4., grad_fn=<PowBackward0>)

tensor(4., grad_fn=<PowBackward0>)

BP

```
1 y.backward()
```

✓ 0.0s

BP

```
1 y.backward()
```

✓ 0.0s

BP

```
1 y.backward()
```

✓ 0.0s

Python

BP

```
1 y.backward()
✓ 0.0s
```

BP

```
1 y.backward()
✓ 0.0s
```

BP

```
1 y.backward()
✓ 0.0s
```

GRAD

```
1 x.grad
✓ 0.0s
tensor(4.)
```

GRAD

```
1 x.grad
✓ 0.0s
tensor(8.)
```

GRAD

```
1 x.grad
✓ 0.0s
tensor(12.)
```

```
1 x = torch.tensor(2.0, requires_grad=True)
2 x
✓ 0.0s
```

tensor(2., requires_grad=True)

```
1 y = x ** 2
2 y
✓ 0.0s
```

tensor(4., grad_fn=<PowBackward0>)

```
1 y.backward()
✓ 0.0s
```

```
1 x.grad
✓ 0.0s
```

tensor(4.)

```
1 x.grad.zero_()
✓ 0.0s
```

tensor(0.)

— x — x — x — x — x — x —

Disabling Gradient Tracking

Sometimes, we may need to perform computations without tracking gradients or calculating derivatives.

Scenarios Where Disabling Gradient Tracking is Useful

- **Model Inference:**
When making predictions with a trained model, gradients are not needed.
- **Model Evaluation:**
During validation or testing phases, to save memory and computation.
- **Feature Extraction:**
When using a model to extract features from data without updating weights.
- **Saving/Loading Model Outputs:**
When storing intermediate results for later use.
- **Visualizations:**
When plotting or analyzing outputs that do not require gradients.
- **Deployment:**
In production environments where only forward passes are performed.

In such cases, PyTorch provides several ways to disable gradient tracking:

- **Set `requires_grad` to `False`:**
You can turn off gradient tracking for a tensor by setting its `requires_grad` attribute to `False` using `requires_grad_(False)`.
- **Detach a tensor:**
Use `.detach()` to create a new tensor that does not require gradients and is disconnected from the computation graph.
- **Use `torch.no_grad()` context:**
Wrap your code inside a `with torch.no_grad():` block to temporarily disable gradient tracking for all operations within the block.

Disabling gradient tracking is useful for inference, evaluation, or any scenario where derivatives are not needed, as it reduces memory usage and speeds up computations.